

Do's

- Define clear business use cases before starting development
- Identify measurable outcomes such as cost reduction, productivity, or accuracy improvement
- Use approved enterprise data sources like SharePoint, databases, or data lakes
- Implement proper access control using Microsoft Entra ID and role-based permissions
- Follow data governance and compliance policies from the organization
- Use Azure AI Search or equivalent for structured retrieval and indexing
- Apply hybrid search (vector + keyword) for better retrieval accuracy
- Design modular architecture separating ingestion, retrieval, and generation layers
- Use evaluation frameworks to measure relevance, groundedness, and hallucination rate
- Implement logging and monitoring for all AI interactions
- Optimize prompts and use query rewriting for better model performance
- Use caching and batching to reduce cost and latency
- Secure API keys, secrets, and tokens using Key Vault or secure environment storage
- Conduct regular testing with real enterprise datasets
- Use version control for prompts, workflows, and pipelines
- Start with pilot projects before scaling enterprise-wide
- Ensure human-in-the-loop for critical or high-risk decisions
- Document architecture, workflows, and data flow clearly
- Train users and stakeholders before deployment
- Continuously improve models using feedback loops and evaluation results

Don'ts

- Do not deploy AI without a defined business problem
- Do not use unapproved or sensitive data without compliance review

- Do not expose API keys, secrets, or credentials in code or logs
- Do not rely on raw LLM output without grounding or retrieval
- Do not ignore access control and data security requirements
- Do not skip evaluation or testing before production deployment
- Do not overload the system with large unchunked documents
- Do not ignore hallucination risks in generative AI outputs
- Do not hardcode prompts without versioning or governance
- Do not build tightly coupled systems that are hard to scale or modify
- Do not skip monitoring and observability setup
- Do not deploy directly to production without pilot validation
- Do not allow unrestricted access to enterprise knowledge bases
- Do not ignore cost implications of large-scale LLM usage
- Do not mix development, testing, and production environments
- Do not assume LLM outputs are always correct or factual
- Do not bypass organizational security policies for convenience
- Do not neglect data privacy regulations such as GDPR or internal policies
- Do not ignore feedback from users and stakeholders
- Do not treat AI systems as static; they require continuous improvement